

Command Broker / Influencer Broker

Qualification and Onboarding Standard

Cross-Sector Framework for Competency Assessment, Role Obligations, and Disclosure Architecture

Implementing Document to the Command Broker Framework, Cross-Sector (CB-Framework_Cross_Sector v0.2)

Version 0.1 — DRAFT

May 2026

Tim Roxey

Eclectic Technologies

For use by organizations deploying Gen AI in consequential domains, sector regulators, compliance officers, legal counsel, and individuals appointed to CB or IB roles

1. Purpose and Scope

This document establishes the qualification and onboarding standard for the Command Broker (CB) and Influencer Broker (IB) roles defined in the Command Broker Framework, Cross-Sector (CB-Framework_Cross_Sector v0.2). It serves two distinct functions that must be read as a unified requirement.

First, it is an onboarding instrument. It states, in explicit terms, what an individual appointed to a CB or IB role is taking on — the obligations, the accountability, and the personal professional consequences of that appointment. An individual who has read and understood this document knows what the role requires of them before they sign anything for the first time.

Second, it is a qualification methodology. It defines the competency domains that must be assessed before an individual may be appointed to a CB or IB role, the assessment criteria against which those competencies are evaluated, and the organizational structure requirements that must be in place for the role to function as designed. Regulators, legal counsel, and compliance officers can assess whether a given individual's appointment satisfies this standard and whether the organizational conditions necessary for the role to be genuine are present.

The standard is organized at the cross-sector level. The qualification methodology — the framework for how competency is assessed — and the onboarding structure — the categories of obligation every CB and IB carries regardless of sector — are established here. Sector-specific content: consequence thresholds, operational context, technical competency bases, and scenario-specific qualification requirements, is carried in the sector annexes to the CB-Framework. This document points to those annexes wherever sector-specific content governs implementation.

Nothing in this document limits the right of any sector, organization, or jurisdiction to impose requirements more stringent than those set out here. This standard establishes a floor, not a ceiling.

2. The Accountability Principle

The foundational requirement of both the CB and IB roles is personal accountability. This requirement is not procedural. It is the load-bearing principle from which every other requirement in this document derives.

A CB authorizing a change to a physical process parameter and an IB authorizing the release of an influence campaign are making decisions with consequential real-world effects. A human being makes those decisions. The accountability for those decisions rests with that human being. It does not rest with the organization that deployed the Gen AI system. It does not rest with the vendor that developed the Gen AI system. It does not rest with the Gen AI system itself.

This principle is not novel. It is the principle established at Nuremberg and affirmed repeatedly in the subsequent development of international law, professional licensing, and regulatory accountability frameworks: that consequential decisions made by human beings in positions of authority carry personal professional and legal consequences for those human beings, and that neither institutional membership nor technological intermediation transfers that accountability elsewhere. There is no chair at Nuremberg for an organization. There is no chair at Nuremberg for a data center. The individual who made the decision, who had the authority and the information to make a different decision, bears the accountability for the decision they made.

Applied to the CB and IB roles: the individual who signs — whose name appears on the authorization record — is accountable for what they authorized. They authorized it after evaluating it and deciding to proceed. That is what their signature means. That is what this standard is designed to ensure: that the evaluation be genuine, that the information be adequate, and that the authority to reject be real.

2.1 No Proxy Authorization

The CB and IB sign in their personal names. No organizational title, role stamp, or delegated authority satisfies the personal name requirement. No individual may sign on behalf of another individual who holds the CB or IB appointment. If the appointed CB or IB is unavailable, a qualified alternate who independently meets this standard may sign — but that alternate signs in their own name, not in the name of the unavailable appointee.

An organization that structures its CB or IB role so that the signing individual has no genuine authority to reject, no genuine access to the required information, or no genuine time to evaluate has not satisfied the accountability requirement. The form is present; the substance is absent. This standard addresses that risk in Section 4 (organizational structure requirements) and Section 6 (IB disclosure architecture). The accountability principle governs: if the signature is not the product of a genuine, informed evaluation by the named individual, it does not provide the protection it purports to provide.

2.2 Personal Signature Requirement

Every authorization produced under this standard must carry the full legal name of the authorizing individual, the date and time of authorization, and a reference to the specific item or campaign authorized. Initials, role titles, electronic tokens not linked to a verified individual identity, and organizational seals do not satisfy this requirement. Where electronic signature systems are used, the signature must be cryptographically linked to the individual's verified identity and the specific document version being authorized.

3. Onboarding: Obligations Common to Both Roles

The following obligations apply to every individual appointed to a CB or IB role, regardless of sector or deployment context. These obligations attach at the moment of appointment and continue for its duration. An individual who accepts a CB or IB appointment accepts these obligations. There is no partial acceptance.

3.1 Obligation to Evaluate Independently

The CB and IB are not relay functions. The role's obligation is not to transmit the Gen AI system's recommendation for execution. It is to evaluate that recommendation against independent knowledge, current conditions, and professional judgment before deciding whether to authorize, modify, or reject. An individual who consistently authorizes without evidence of independent evaluation is not performing the role. This obligation requires that the individual have genuine competency in the subject matter of their evaluations — competency sufficient to recognize when a Gen AI recommendation is inconsistent with conditions, constraints, or applicable requirements, not merely to read and acknowledge what the recommendation says.

3.2 Obligation to Reject

The authority to authorize is only meaningful if the authority to reject is equally real. A CB or IB who has never rejected a Gen AI recommendation is a candidate for inquiry, not commendation. The organizational structure surrounding the CB and IB roles must make rejection operationally feasible: no workflow design, performance metric, throughput expectation, or organizational pressure may be structured to make rejection practically difficult or professionally costly. An individual who believes that their authority to reject is being structurally undermined has an affirmative obligation to report that condition to the appropriate oversight authority and, if necessary, to decline to sign.

3.3 Obligation of Situational Awareness

The CB and IB must have real-time or contemporaneous access to the information necessary to evaluate what they are authorizing. A CB who cannot see the process data underlying a Gen AI recommendation is not in a position to evaluate the recommendation independently. An IB who cannot access the actual content of a campaign they are authorizing cannot verify that the system-generated disclosure form accurately characterizes it. The obligation of situational awareness requires the CB and IB to affirmatively verify, before signing, that they have sufficient information to do so. If they do not have sufficient information, they do not sign until they do.

3.4 Obligation of Competency Maintenance

Appointment to a CB or IB role is not a permanent qualification. The processes being managed change. The Gen AI systems being reviewed change. The threat environment changes. A CB or

IB whose competency was sufficient at the time of appointment but who has not maintained currency with the relevant operational environment, the applicable technical domain, or the evolving capabilities and failure modes of the Gen AI systems they are reviewing is not performing the role to the standard required. Each appointing organization must establish a competency maintenance schedule appropriate to the sector annex requirements. That schedule and evidence of compliance with it are part of the audit record.

3.5 Obligation of Record Integrity

The authorization record produced by a CB or IB is a legal document. It must accurately reflect what was reviewed, when it was reviewed, what information was available at the time of review, and what was authorized. An individual who signs an authorization record that does not accurately reflect those elements — whether by commission, omission, or acquiescence to a pre-filled record they did not verify — has produced a false record. The obligation of record integrity requires that the CB or IB verify the accuracy of the authorization record before signing it, rather than merely signing what is placed before them.

3.6 Obligation of Stress-Condition Performance

The CB and IB roles are most consequential when conditions are most difficult. A CB may be called upon to evaluate Gen AI recommendations under degraded sensor feeds, loss of communications, elevated operational tempo, or conditions arising from a cyberattack or military action. An IB may be called upon to authorize content under time pressure, with incomplete information, or in a rapidly evolving operational environment. The obligation of stress-condition performance means that the individual appointed to these roles must be qualified to perform them under adverse conditions, not only under normal operating assumptions. The qualification assessment must include an evaluation under simulated stress conditions as defined in the applicable sector annex.

4. Qualification Methodology

The qualification methodology establishes how competency for CB and IB roles is assessed. It defines the competency domains that must be evaluated, the assessment approach for each domain, and the organizational structure requirements that must be verified independently of the individual candidate's competency. Sector annexes provide domain-specific content within this structure.

Qualification is not a one-time event. It is a determination, made at appointment and periodically thereafter, that a named individual meets the competency requirements for a defined role in a defined operational context. The qualification record must identify: the individual assessed, the role and context, the date of assessment, the domains assessed, the assessment methodology

used, the outcome of each domain assessment, and the name and qualification of the assessor. That record is an audit artifact and must be preserved for the duration of the individual's appointment plus a period defined in the applicable sector annex.

4.1 Competency Domains — Command Broker

The following domains must be assessed for every CB candidate. The applicable sector annex establishes the assessment criteria within each domain.

Technical process competency. The CB must demonstrate sufficient knowledge of the physical process within their authorization scope to evaluate Gen AI recommendations independently. Assessment must verify that the candidate can: identify when a Gen AI recommendation is inconsistent with current process conditions; recognize parameter values or recommended changes that approach or exceed operational limits; and distinguish between normal process variation and anomalous conditions that require a response different from the Gen AI recommendation. Assessment methodology: scenario-based evaluation using representative process conditions, including at least one scenario in which the Gen AI recommendation is incorrect and one in which it is correct but requires contextual judgment to confirm.

Authority and role clarity. The CB must demonstrate that they understand the scope of their authorization authority: what they may authorize, what falls outside their scope, and which escalation path applies when a recommendation does. Assessment methodology: structured interview or written examination covering authorization scope, escalation protocol, and the conditions under which sign-off must be withheld.

Situational awareness under normal conditions. The CB must demonstrate the ability to integrate Gen AI advisory output with direct observation of process conditions and available sensor data to reach an independent authorization decision. Assessment methodology: simulation exercise with access to process data alongside Gen AI recommendations; evaluator assesses whether the candidate's decision-making reflects integration of both sources or reliance on Gen AI output alone.

Stress-condition performance. The CB must demonstrate the ability to perform the broker function under degraded conditions as defined in the applicable sector annex. Minimum stress scenarios for assessment: loss of one or more sensor feeds; communications degradation requiring local-only decision-making; and elevated tempo requiring evaluation of multiple sequential recommendations within a compressed timeframe. Assessment methodology: scenario-based simulation under each stress category, with evaluator assessment of decision quality and process.

Adversarial scenario awareness. The CB must demonstrate awareness of the conditions under which a Gen AI recommendation may itself be the product of adversarial manipulation — prompt injection, training data compromise, or sensor feed falsification. The CB does not need to be a cybersecurity analyst; they need to recognize indicators that a recommendation is

anomalous, suggesting manipulation rather than normal process variation. Assessment methodology: scenario-based evaluation, including at least one scenario in which the Gen AI recommendation reflects adversarially manipulated inputs.

4.2 Competency Domains — Influencer Broker

The IB qualification domains reflect the distinct cognitive demands of the role. Two authorization levels are defined: content-level authorization (a single content item or a small, defined set) and campaign-level authorization (a distributed content architecture that reaches a defined audience at scale). Individuals may be qualified at one or both levels. Campaign-level qualification requires content-level qualification as a prerequisite.

4.2.1 Content-Level IB Qualification Domains

Editorial judgment. The IB must demonstrate the ability to evaluate AI-generated content for accuracy, contextual appropriateness, internal consistency, and alignment with the generating organization's declared intent. Assessment methodology: structured review of a set of AI-generated content samples, some accurate and appropriately framed, some inaccurate, some accurate but misleading through selective framing. The evaluator assesses whether the candidate identifies each category correctly and articulates the basis for their assessment.

Source and claim verification. The IB must demonstrate the ability to identify factual claims in AI-generated content and assess whether the sourcing provided in the disclosure form is adequate to support those claims. The IB is not required to verify every claim independently; they are required to assess whether the disclosure form identifies the claims and whether the sourcing is plausible and traceable. Assessment methodology: review of disclosure forms accompanying sample content items, with evaluator assessment of whether the candidate identifies unsourced or inadequately sourced claims.

Risk parameter evaluation. The IB must demonstrate the ability to evaluate risk flags identified in the disclosure form and assess whether the risk characterization is accurate and complete. The IB must also be capable of identifying risk parameters that the disclosure form has not flagged. Assessment methodology: review of disclosure forms with intentionally incomplete risk flag inventories; evaluator assesses whether the candidate identifies the missing flags.

Disclosure form verification. The IB must demonstrate understanding of the disclosure form architecture, including the diversity requirement for the disclosure system (Section 6.3), and the ability to verify that the form was generated by a system meeting that requirement. Assessment methodology: structured examination on disclosure form requirements, including the diversity requirement and the conditions under which a form is insufficient to support authorization.

4.2.2 Campaign-Level IB Qualification Domains

Campaign-level qualification requires all content-level domains above, plus the following additional domains.

Campaign architecture evaluation. The campaign-level IB must demonstrate the ability to evaluate a distributed content architecture: targeting logic, distribution channel selection, audience definition, and timing trigger design. The IB is not authorizing individual content instances; they are authorizing the architecture within which those instances will be produced and distributed. Assessment must verify that the candidate understands what they are and are not authorizing when they sign a campaign-level disclosure form.

Temporal validity judgment. The campaign-level IB must demonstrate the ability to assess whether the conditions assumed in a campaign authorization remain valid at execution, and to determine when changed conditions require re-authorization before the campaign proceeds. Assessment methodology: scenario-based evaluation in which an authorized campaign encounters changed factual conditions between authorization and execution; the evaluator assesses whether the candidate correctly identifies the re-authorization trigger.

Scale and consequence assessment. The campaign-level IB must demonstrate the ability to assess the potential consequences of a campaign that distributes to a defined audience at the declared scale, including second-order effects that may not be visible at the individual content instance level. Assessment methodology: structured scenario analysis requiring the candidate to identify potential consequences of described campaign architectures at varying distribution scales.

4.3 Organizational Structure Requirements

Individual competency is necessary but not sufficient. The organizational conditions in which the CB or IB operates must also be verified. The following requirements apply to the appointing organization and must be documented in the qualification record.

- Independence from the generating function. The CB or IB must not report to, be evaluated by, or have their compensation determined by the same organizational unit responsible for producing the Gen AI outputs they are authorizing. Structural independence is required. Informal independence is insufficient.
- Protected authority to reject. The organization must document, in writing, that the CB or IB has the authority to reject any Gen AI recommendation or authorization request without organizational sanction. That documentation must be provided to the CB or IB at the appointment, included in the qualification record, and accessible to any oversight authority reviewing the appointment.
- Adequate time allocation. The organization must demonstrate that the CB or IB has been allocated sufficient time to evaluate the volume of authorization requests they will receive independently. An IB who is allocated two hours per week but expected to

process 50 campaign authorization requests per week has not been given adequate time, regardless of the organizational chart.

- Access to required information. The organization must verify that the CB or IB has direct access to the process data, sensor feeds, content samples, and disclosure form documentation required to perform independent evaluations. Intermediated access — where the CB or IB receives summaries prepared by the generating unit — does not satisfy this requirement.
- Qualified alternate. The organization must designate and qualify at least one alternate CB or IB for each role to ensure continuity of the broker function when the primary appointee is unavailable. The alternate must meet the same qualification standard as the primary. Their qualification record must be maintained separately.

5. CB-Specific Onboarding Content

This section is addressed to the individual being appointed as a Command Broker. It states, in direct terms, what you are taking on.

5.1 What You Are Authorizing

When you confirm a Gen AI recommendation as Command Broker, you are authorizing a change to a physical process parameter. That change will propagate to the equipment. If the change is wrong — because the Gen AI recommendation was erroneous, because the conditions under which it was generated have changed, because the recommendation reflected adversarially manipulated inputs, or because you failed to evaluate it adequately — the physical consequence is yours to account for.

You are not the last line of defense in a system with no other protections. The DBA sector framework establishes multiple layers of protection before a consequence becomes irreversible. But you are a defined, essential layer. The architecture was designed with your evaluation as a required element. If you treat the role as a formality, the architecture fails at the layer it was designed to have.

5.2 What Your Signature Means

Your signature on a CB authorization record means that you reviewed the Gen AI recommendation against the available process data. I evaluated the recommendation against my independent knowledge of the process. I determined that the recommended action is consistent with current conditions and applicable operational limits. I authorized the action.

It does not mean "I read the recommendation." It does not mean the system recommended it; I passed it along. It does not mean: my organization instructed me to approve. The signature attests to the evaluation, not merely to the receipt.

5.3 When You Must Not Sign

You must not sign when: you do not have access to the process data needed to evaluate the recommendation independently; the recommendation is anomalous in ways you cannot account for; you are under time pressure that prevents genuine evaluation; you have reason to believe the Gen AI inputs may have been compromised; or the recommended action falls outside your defined authorization scope. In each of these cases, the correct action is to withhold authorization and escalate in accordance with the protocol defined in your sector annex. Withholding authorization when you should not sign is the role working as designed. It is not a failure.

6. IB-Specific Onboarding Content and Disclosure Architecture

This section is addressed to the individual being appointed as an Influencer Broker and to the organizations that deploy Gen AI influence systems subject to this standard. It establishes both the personal obligations of the IB role and the architectural requirements for the disclosure system that supports it.

6.1 What You Are Authorizing

When you sign an IB authorization, you are authorizing the release of AI-generated content into an information environment. At the content level, you are authorizing a specific item or a defined set of items. At the campaign level, you are authorizing an architecture — a defined targeting logic, distribution scope, channel selection, and timing structure — within which content will be produced and distributed at scale.

The information environment does not have the physical reversibility constraints of a process control system, but it has its own irreversibility: content released at scale cannot be recalled. An influence campaign that reaches its intended audience has already had its effect by the time any error in it is identified. The IB's evaluation before release is the only meaningful opportunity for human judgment to be applied. After the campaign is executed, the IB's name is on the authorization record. That record survives.

6.2 What Your Signature Means

Your signature on an IB authorization means that you received the system-generated disclosure form for this content item or campaign. I verified that the disclosure form was generated by a system diverse from the payload-generating system, as required by Section 6.3 of this standard. I reviewed the disclosure form and, at my discretion, reviewed sample content directly. I evaluated the declared intent, payload characterization, distribution scope, risk parameter inventory, and authorization scope. I determined that the content or campaign as characterized is consistent with

the declared intent, that the risk parameters are accurately and completely identified, and that the authorization scope is precisely defined. I authorized the release.

The same principle applies as in the CB role: the signature attests to the evaluation, not merely to the receipt of the form.

6.3 The Disclosure Form Architecture

6.3.1 System-Generated Disclosure Requirement

No IB authorization may proceed without a system-generated disclosure form. The form is not optional, not a courtesy document, and not a summary prepared by the content team. It is a required architectural element — a precondition for release, enforced by the system architecture — such that the release pathway is unavailable until a disclosure form has been generated and routed to the IB.

The disclosure system generates the form using a structured template that specifies the required fields. The deploying organization provides the template. The disclosure system populates the template by reviewing the payload and campaign architecture produced by the payload system. The template must be passed to the disclosure system in blank form — the payload system must not pre-populate template fields before handoff to the disclosure system.

6.3.2 The Diversity Requirement

The system that generates the IB disclosure form must be a different Gen AI system from the system that generated the payload and distribution strategy. This is a diversity requirement, derived from the same principle that governs the design of safety-critical systems in the nuclear sector.

In nuclear facility design, redundancy alone — multiple trains of the same equipment type — protects against random independent failures but not against common-cause failures. A design defect or systematic vulnerability that affects one train of a given pump type affects all trains of that type simultaneously. Diversity is the engineering response: at least one train must use a different technology, so that no single failure mechanism can defeat all trains at once. The turbine-driven pump fails for reasons different from those of the motor-driven pumps. A failure mode that defeats one does not defeat the other.

The same logic applies here. If the payload-generating system and the disclosure-form-generating system are the same model, or models from the same developer with the same training lineage, then whatever causes the payload system to mischaracterize, selectively frame, or omit material parameters will cause the form system to do the same. The IB reviews a document produced by the same failure mode that it is supposed to catch—redundancy without diversity. The diverse disclosure system breaks that common-cause failure mode: it characterizes what it sees in the payload, not what the payload system intended to represent.

The diversity requirement is satisfied when the disclosure system is: a different Gen AI model from the payload system, developed by a different organization. Both systems must be identified by name and version in the disclosure form itself, so the diversity basis is self-documenting in the audit record.

Example of a satisfied diversity requirement: the payload and distribution strategy are developed using one commercially available large language model; the disclosure form is generated using a different large language model from a different developer. Both systems are identified by name and version on the form. The IB verifies this identification before signing.

The diversity requirement does not mandate a certification process, regulatory pre-approval, or independent audit of the diversity basis. It requires documentation. If that documentation is later found to be false — if the same system generated both the payload and the form — that is evidence of deliberate fraud, which is a more serious legal exposure than the underlying campaign would have created.

6.3.3 Minimum Parameter Set — Content-Level Disclosure

A content-level disclosure form must include the following fields, each populated by the disclosure system:

- Content identity. Title or identifier, format, publication channel, and intended audience.
- Generating systems. Name and version of the payload-generating system. Name and version of the disclosure-generating system. Confirmation that the two systems are different, developed by different organizations.
- Intent declaration. The stated purpose of the content, declared by the deploying organization at the time of generation, is specific enough to be auditable against the content itself.
- Factual claim inventory. A complete list of factual claims made in the content, with the sourcing identified by the disclosure system for each claim. Claims for which no sourcing is identified must be flagged explicitly.
- Risk parameter inventory. All elevated risk parameters identified by the disclosure system, including: content about identifiable individuals; electoral or political subject matter; contested factual claims; content that could be mistaken for official communication; synthetic media elements. For each identified parameter, the disclosure system must state the basis for the flag.
- Authorization scope. A precise statement of what the IB is authorizing: this specific content item, in this format, for this publication channel, at this time. Republication, reformatting, or redistribution to a different audience requires a new disclosure form and a new IB authorization.
- IB attestation block. Full legal name field, date and time field, and a structured attestation statement that the IB has reviewed the form, verified the diversity basis, evaluated each parameter, and authorizes the release as scoped.

6.3.4 Minimum Parameter Set — Campaign-Level Disclosure

A campaign-level disclosure form must include all content-level fields above, adapted to campaign scope, plus the following additional fields:

- Campaign identity. Campaign name, unique identifier, generating organization, and authorized IB.
- Payload characterization. Content types included in the campaign, formats, tone categories, and the range of factual claims that may appear across campaign content instances.
- Distribution architecture. Target population definition and selection methodology; distribution channels; order-of-magnitude scale (minimum: the approximate number of individuals who will receive campaign content); timing triggers and the conditions under which each trigger executes.
- Sample content. A structured sample of campaign content instances, selected by the disclosure system using a defined sampling methodology. The sample must be sufficient for the IB to verify that the content is consistent with the payload characterization. The sampling methodology must be specified in the form and must be resistant to selection bias by the payload system.
- Temporal validity parameters. The time window within which this authorization is valid. The factual and contextual conditions under which the authorization holds. The conditions that would require re-authorization before campaign execution if they change after the authorization date.
- Campaign-level risk assessment. An assessment of potential second-order effects at the declared distribution scale, not only at the individual content instance level.
- Authorization scope statement. An explicit statement of what the IB is and is not authorizing: the campaign architecture as described, within the defined time window, under the stated conditions. Individual content instances produced within the authorized architecture are covered by this authorization only if they are consistent with the reviewed payload characterization. Instances outside that characterization require a separate authorization.

6.3.5 Temporal Validity and Re-Authorization

A campaign authorization is not a permanent authorization. It is valid for the time window specified in the disclosure form, under the conditions stated therein. If the campaign has not been executed before the authorization window expires, re-authorization is required. If factual or contextual conditions specified as holding conditions change materially before execution, the system must flag the change, and re-authorization is required before execution proceeds.

The IB who reauthorizes a campaign following a change in condition is making a new authorization determination. The re-authorization is a separate signature event that generates a separate audit record, documenting the changed condition and the basis for the new authorization.

6.4 When the IB Must Not Sign

The IB must not sign when: the disclosure form was not generated by a system diverse from the payload system; the form does not contain the required minimum parameters; the form was pre-populated by the payload system before handoff to the disclosure system; the IB does not have access to the sample content required to verify the payload characterization; the factual claims in the content are not adequately sourced; the risk parameter inventory is incomplete; the authorization scope is insufficiently defined to support a meaningful authorization; or the IB is under time or organizational pressure that prevents genuine evaluation.

In each of these cases, the IB returns the authorization request with a documented statement of the specific deficiency. The statement is an audit artifact. The deploying organization is responsible for remedying the deficiency before the IB is asked to sign again.

7. Audit Record Requirements

Every CB and IB authorization produces an audit record. The audit record is a legal document. The following elements must be present in every audit record:

- The complete authorization form as it existed at the moment of signature, including all fields populated by the disclosure system (IB) or the operational system (CB).
- The full legal name of the authorizing individual and the date and time of authorization.
- For IB authorizations: the name and version of both the payload system and the disclosure system, confirming the diversity basis.
- For campaign-level IB authorizations: the sample content reviewed, the sampling methodology, and the temporal validity parameters.
- Any conditions or limitations noted by the CB or IB at the time of authorization.
- Any re-authorization events, with documentation of the changed conditions that triggered re-authorization.
- The qualification record of the authorizing individual, identifying the assessment date, the domains assessed, and the assessor.

Audit records must be preserved for a minimum period defined in the applicable sector annex. In the absence of a sector annex specification, the minimum preservation period is seven years from the date of authorization or the date of the last campaign execution under that authorization, whichever is later. Records must be stored in a format that is retrievable and readable for the duration of the preservation period, independent of changes to the systems that generated them.

8. Relationship to Sector Annexes

This standard establishes the cross-sector qualification methodology and onboarding structure. Sector annexes to the CB-Framework define the sector-specific content that populates this structure for each deployment context. Where a sector annex provision conflicts with this standard, the more stringent requirement governs.

Sector annexes currently in development or anticipated: DBA-W (water sector), DBA-GAS (natural gas transmission), DBA-OIL (petroleum sector), DBA-Nuclear (nuclear facilities), and the Democratic Safeguards project (information environment). Organizations operating in sectors without a developed annex must apply this standard directly, using consequence analysis appropriate to their operational context to determine tier assignments, competency baselines, and stress scenario content.

9. Foundational References

CB-Framework_Cross_Sector v0.2. Roxey, T. Eclectic Technologies, 2026. [Parent document establishing the broker concept, the Bright Line, and the three-tier criticality framework.]

CB-CT_Commissioning_and_Joint_Training v0.2. Roxey, T.E. Eclectic Technologies, 2026. [Companion module sequencing commissioning and joint training of the Command Broker at the vendor–AOO seam.]

Roxey, T.E. Generative AI Risk in Industrial Control Systems: A Study for Critical Infrastructure Operators, Regulators, and Suppliers. Version 29. Eclectic Technologies, 2026. [Primary source for the three-tier criticality framework.]

DBA-W-Sector_Framework_v1_0. Roxey, T. Eclectic Technologies. [Water sector Design Basis. Section 5.5 establishes the command-broker architecture as a sector Design Basis requirement.]

Democratic Safeguards Project. Roxey, T. Eclectic Technologies. [In development. Source for the Influencer Broker concept and democratic governance application.]

Revision History

Version	Date	Description
0.1	May 2026	Initial draft. Establishes cross-sector qualification methodology,

Version	Date	Description
		onboarding obligations common to both roles, CB-specific and IB-specific onboarding content, IB disclosure form architecture, including the diversity requirement and minimum parameter sets at content and campaign level, audit record requirements, and relationship to sector annexes.
0.1 (housekeeping)	July 2026	Housekeeping correction, two items, no substantive content change. (1) “Brightline” corrected to “Bright Line” in the References section, per the series-wide convention in DBA-EN-SF §4.3. (2) Source-study citation updated to the canonical edition — Roxey, T.E., “Generative AI Risk in Industrial Control Systems,” Version 29, Eclectic Technologies (2026) — replacing the prior ISAGCA-attributed “A Study of AI Risks — Vendor Agnostic,” Version 28.
0.1 (cite refresh)	Jul 2026	Citations refreshed to current corpus versions: CB-Framework v0.2, Command Broker Implementation Guide V0.2, CB-CT v0.2, DBA-VC v0.8. Mechanical version-reference pass; no substantive content change. (Held pending decision: the NERC/BES-Variant TAF reference and the ISA Parent Framework V1.2 label.)